

Portland State University

Project Report

Title: Hardware Acceleration of Genetic Algorithm Mutation Step

Student: Bhavana Manikyanahalli Srinivasegowda

Objective

The goal of this project is to accelerate the mutation step of a Genetic Algorithm (GA) used for string matching tasks by implementing it in hardware. This project focuses on improving performance and energy efficiency by offloading mutation from Python-based software to a synthesized hardware module.

Background

Traditional GAs are entirely implemented in software. While flexible, software-only implementations are relatively slow and power-hungry, especially in edge devices or embedded systems. Among the core GA steps—fitness, selection, crossover, and mutation—the mutation step is simple, highly parallel, and well-suited for hardware acceleration.

Methodology

- Implemented the mutation step in SystemVerilog, supporting configurable gene length and mutation rate.
- Verified the design using Cocotb and Python testbenches.
- Synthesized and evaluated the hardware using OpenLane and analyzed power, timing, and area metrics.
- Compared performance against a baseline Python implementation.

Results

Technology & Flow Info:

- Toolchain: OpenLane
- Technology Node: Sky130
- Flow Used: Classic
- Run Tag: RUN_2025-06-11_18-46-38

Area & Utilization:

- Standard Cell Count: 83,330
- Fill Cells: 493,410
- Core Instance Area: 114,014 μm^2
- Die Area: 9,000,000 nm^2 (3 mm × 3 mm)
- Core Area: 5.75 mm^2
- Utilization: ~1.98%

Power:

- Total Power: 3.96 μ W
 - Internal: 1.79 μ W
 - Switching: 2.17 μ W
 - Leakage: \sim 0.34 nW

Timing:

- Worst Setup Slack (Best Corner): +5.87 ns (MET)
- Worst Hold Slack (Best Corner): +0.13 ns (MET)
- Setup Violations (Worst Corner): 138
- Hold Violations (Worst Corner): 4
- Max Skew: 0.36 ns
- Estimated Maximum Frequency: \sim 170 MHz
 - (based on worst setup time)

Signoff Checks:

- Max Fanout Violations: 3
- Max Slew Violations: 0
- Max Cap Violations: 0
- Unannotated Drivers: 8
- Routing DRC Errors: 0
- Magic DRC Errors: 0
- LVS Errors: 0
- Antenna Violations: 3

Wirelength & IR Drop:

- Total Wirelength: 212,214 units
- IR Drop (Worst): 0.62 mV
- Lowest Observed Voltage: 1.799 V

Mutation Analysis:

- Mutation Types Tested: 6

Fitness function minimized error between expected and generated results using mean squared error.

CProfiling shows execution time bottleneck reduced after optimizing mutation logic.

```

MODULE=test_mutation_sv TESTCASE= TOPLEVEL=mutation_sv TOPLEVEL_LANG=verilog \
sim_build/Vtop
-.-ns INFO gpi .mbed/gpi_embed.cpp:108 in set_program_name_in_venv Using Python vir
tual environment interpreter at /home/sirim/week8/cocotb-env/bin/python
-.-ns INFO gpi .:/gpi/GpiCommon.cpp:101 in gpi_print_registered_impl VPI registered
0.00ns INFO cocotb Running on Verilator version 5.020 2024-01-01
0.00ns INFO cocotb Running tests with cocotb v1.9.2 from /home/sirim/week8/cocotb-env/lib/python
3.12/site-packages/cocotb
0.00ns INFO cocotb Seeding Python random module with 1748926348
0.00ns INFO cocotb.regression pytest not found, install it to enable better AssertionError messages
0.00ns INFO cocotb.regression Found test test_mutation_sv.test_mutation_sv
0.00ns INFO cocotb.regression running test_mutation_sv (1/1)
Original gene : Hello World!
Mutated gene : Hello {orl|
155.00ns INFO cocotb.regression test_mutation_sv passed
155.00ns INFO cocotb.regression *****
***** Target Word : Hello World! *****
ATIO (ns/s) Max Population : 10
***** Mutation Rate : 0.2 *****
57385.59 -----
***** The Best Fitness Time *****
17449.17 -----
*****
D?fTH v""OPD 8.33% 0:00:00.546480
S?fTH o0;l2^ 16.67% 0:00:00.563878
D?fTo o0;l2^ 25.0% 0:00:00.574289
y?fTo o0rl2^ 33.33% 0:00:00.610398
D?fTo ourl2! 41.67% 0:00:00.650839
yefTo ourl2! 50.0% 0:00:00.661524
TelTo o{rl2! 58.33% 0:00:00.713715
Tel(o oorl2! 66.67% 0:00:00.743458

```

Fig1: Cocotb results

Fig2: Python results

```
221826 function calls (221646 primitive calls) in 6.816 seconds
```

Ordered by: internal time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
9810	5.507	0.001	5.963	0.001	GA1.py:47(mutation)
39250	0.505	0.000	0.505	0.000	GA1.py:12(calculate_fitness)
9810	0.261	0.000	0.512	0.000	GA1.py:38(crossover)
1	0.121	0.121	6.816	6.816	GA1.py:79(main)
19852	0.092	0.000	0.092	0.000	{method 'join' of 'str' objects}
47402	0.071	0.000	0.071	0.000	{built-in method builtins.chr}
19697	0.053	0.000	0.053	0.000	{built-in method builtins.max}
19643	0.041	0.000	0.094	0.000	GA1.py:69(bestfitness)
10/6	0.024	0.002	0.087	0.015	{built-in method _imp.exec_dynamic}
9822	0.024	0.000	0.024	0.000	{built-in method builtins.round}
10	0.022	0.002	0.022	0.002	{built-in method _imp.create_dynamic}
39819/39817	0.017	0.000	0.017	0.000	{built-in method builtins.len}

Fig3: C Profiling Output

Conclusion

The project demonstrates that accelerating only the mutation step can provide significant speed and energy gains. The hardware module achieved over 8 million mutations per second at just 3.96 μ W power, validating the benefit of selective hardware offloading in genetic algorithms.